

# SVM Week 1

*David Milmont*

*May 21, 2017*

## Loading Packages and Downloading Data - Creating Train, Test, Validate data sets

```
library('kernlab')
library('RCurl')
```

```
## Loading required package: bitops
```

```
library('ggplot2')
```

```
##
```

```
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:kernlab':
```

```
##
```

```
##      alpha
```

```
library('GGally')
library('mlr')
```

```
## Loading required package: ParamHelpers
```

```
library('kknn')
```

```
file <- getURL('https://d37djvu3ytnwxt.cloudfront.net/assets/courseware/v1/e39a3df780dacd5503df6a8322d7')
```

```
data <- read.csv(textConnection(file), header=T, sep = "\t")
```

```
#Summary of data
```

```
summarizeColumns(data)
```

```
##      name      type na      mean      disp  median      mad      min
## 1      A1 integer  0    0.6896024  0.4630105    1.000    0.000000    0.00
## 2      A2 numeric  0   31.5783486  11.9817891   28.460   10.318896   13.75
## 3      A3 numeric  0    4.8305581    5.0232952    2.855    3.335850    0.00
## 4      A8 numeric  0    2.2416896    3.3691972    1.000    1.356579    0.00
## 5      A9 integer  0    0.5351682    0.4991434    1.000    0.000000    0.00
## 6     A10 integer  0    0.5611621    0.4966249    1.000    0.000000    0.00
## 7     A11 integer  0    2.4984709    4.9656552    0.000    0.000000    0.00
## 8     A12 integer  0    0.5382263    0.4989182    1.000    0.000000    0.00
## 9     A14 integer  0  180.0840979  168.3157190  160.000  148.260000    0.00
## 10    A15 integer  0 1012.7308869  5249.3206597    5.000    7.413000    0.00
## 11     R1 integer  0    0.4525994    0.4981291    0.000    0.000000    0.00
##              max nlevs
## 1           1.00      0
## 2          80.25      0
## 3          28.00      0
## 4          28.50      0
## 5           1.00      0
```

```
## 6      1.00      0
## 7     67.00      0
## 8      1.00      0
## 9    2000.00      0
## 10 100000.00      0
## 11      1.00      0

#Create data frame split
set.seed(546)

#shuffling to ensure randomness
data <- data[sample(nrow(data)),]

#Getting idea of sizes
nrow(data) * .60

## [1] 392.4
nrow(data) * .20

## [1] 130.8
nrow(data) * .20

## [1] 130.8

#Splitting Manually - KISS method
train <- data[1:394,]
test <- data[395:525,]
validate <- data[526:654,]
```

## Visualization

Quick visual to see how the data is correlated

```
data[,11] <- as.factor(data[,11])
GGally::ggpairs(data[, c(2:4,9:11)], aes(colour=R1))
```



## Creating the Model - SVM

Selected the polydot kernel as it has the highest accuracy. After running MLR package for parameter hypertuning settled with  $C=2.73e+06$  for 100% accuracy - This results in a very small margin hyperplane. I suspect in a real life scenario this is not optimal as the use case would benefit from some margin of error for review scenarios.

```
#Creating matrix to train model
x <- as.matrix(data[,1:10])
#creating target
y <- data[,11]

model <- ksvm(y ~ x, type = "C-svc", kernel = "polydot", C=2.73e+06, scaled = TRUE, cross = 5)
```

```

a <- colSums(data[model@SVindex,1:10]*model@coef[[1]])
cat('a:',a,'\n')

## a: -298.9353 -16576.53 -4488.99 2362.042 536.0985 -382.8606 2740.088 -491.4585 8406.12 4124302
a0 <- sum(a*data[1,1:10]) - model@b
cat('a0:',a0,'\n')

## a0: 4127211276
pred <- predict(model,data[,1:10])
# pred
data$prediction <- pred

cat('SVM accuracy:',sum(pred == data[,11]) / nrow(data),'\n')

## SVM accuracy: 1
cat('offset:',b(model),'\n')

## offset: 8.856835
cat('error',error(model),'\n')

## error 0
kernelnf(model)

## Gaussian Radial Basis kernel function.
## Hyperparameter : sigma = 0.0916808764420663

```

## Kernel Selection

Checking different kernels and their affect on prediction accuracy - selected polydot for final model

```

kernels <- c('rbfdot','polydot','vanilladot','tanhdot','laplacedot','besseldot','anovadot','splinedot',
             'stringdot')

for(kernel in kernels){
  model <- ksvm(x, y, type = "C-svc", kernel = kernel, C=100, scaled = TRUE, cross = 5)
  pred <- predict(model,data[,1:10])
  cat('\n',kernel,'pred: ', sum(pred == data[,11]) / nrow(data))
}

##
## rbfdot pred: 0.9587156
## polydot pred: 0.9525994
## vanilladot pred: 0.9541284
## tanhdot pred: 0.9541284
## laplacedot pred: 0.9587156
## besseldot pred: 0.9541284
## anovadot pred: 0.9480122
## splinedot pred: 0.9541284
## stringdot pred: 0.9571865

```

## Paramater Hypertuning

Utilized the MLR package for parameter tuning, selected C=2.73e+06 for highest prediction accuracy

*#Trying mlr package for parameter hypertuning*

```
trainTask <- makeClassifTask(data = data, target = 'R1')
trainTask
```

```
## Supervised task: data
## Type: classif
## Target: R1
## Observations: 654
## Features:
## numerics  factors  ordered
##      10      1      0
## Missings: FALSE
## Has weights: FALSE
## Has blocking: FALSE
## Classes: 2
##   0   1
## 358 296
## Positive class: 0
```

```
learner <- makeLearner("classif.ksvm")
```

```
ksvm <- makeLearner("classif.ksvm", predict.type = "response")
getParamSet("classif.ksvm")
```

##	Type	len	Def		Constr	Req	Tunable	Trafo
## scaled	logical	-	TRUE		-	-	TRUE	-
## type	discrete	-	C-svc		-	-	TRUE	-
## kernel	discrete	-	rbfdot		-	-	TRUE	-
## C	numeric	-	1		0 to Inf	Y	TRUE	-
## nu	numeric	-	0.2		0 to Inf	Y	TRUE	-
## epsilon	numeric	-	0.1		-Inf to Inf	Y	TRUE	-
## sigma	numeric	-	-		0 to Inf	Y	TRUE	-
## degree	integer	-	3		1 to Inf	Y	TRUE	-
## scale	numeric	-	1		0 to Inf	Y	TRUE	-
## offset	numeric	-	1		-Inf to Inf	Y	TRUE	-
## order	integer	-	1		-Inf to Inf	Y	TRUE	-
## tol	numeric	-	0.001		0 to Inf	-	TRUE	-
## shrinking	logical	-	TRUE		-	-	TRUE	-
## class.weights	numericvector	<NA>	-		-	-	TRUE	-
## fit	logical	-	TRUE		-	-	TRUE	-
## cache	integer	-	40		-	-	TRUE	-
##								
## scaled					-	-	TRUE	-
## type	C-svc,nu-svc,C-bsvc,spoc-svc,kbb-svc				-	-	TRUE	-
## kernel	vanilladot,polydot,rbfdot,tanhdot,lap...				-	-	TRUE	-
## C					0 to Inf	Y	TRUE	-
## nu					0 to Inf	Y	TRUE	-
## epsilon					-Inf to Inf	Y	TRUE	-
## sigma					0 to Inf	Y	TRUE	-
## degree					1 to Inf	Y	TRUE	-
## scale					0 to Inf	Y	TRUE	-
## offset					-Inf to Inf	Y	TRUE	-
## order					-Inf to Inf	Y	TRUE	-
## tol					0 to Inf	-	TRUE	-

```
## shrinking - - TRUE -
## class.weights 0 to Inf - TRUE -
## fit - - FALSE -
## cache 1 to Inf - TRUE -

set_cv <- makeResampleDesc("CV",iters = 3L)
pssvm <- makeParamSet(
  makeNumericParam("C", lower = -10, upper = 10, trafo = function(x) 10^x),
  makeNumericParam("sigma", lower = -10, upper = 10, trafo = function(x) 10^x)
)

ctrl = makeTuneControlRandom(maxit = 200L)

res <- tuneParams(ksvm, task = trainTask, resampling = set_cv, par.set = pssvm, control = ctrl)

t.svm <- setHyperPars(ksvm, par.vals = res$x)

par.svm <- train(ksvm, trainTask)
predict.svm <- predict(par.svm, trainTask)

res

## Tune result:
## Op. pars: C=8.36e+07; sigma=0.00261
## mmce.test.mean= 0
```

**Training KKNn - Using Test, Train, And Validate data sets - results are KKNn accuracy: 0.9160305**

```
train.knn <- kknn(formula = formula(train$R1~.), train = train, test = test, k = 7, distance = 1, kernel = "euclidean")

fit <- fitted(train.knn)

table(test$R1, fit)

##      fit
##      0 0.0127439975524816 0.0405412744063402 0.0727517584925532
## 0 38                2                3                1
## 1 4                 0                0                0
##      fit
##      0.0854957560450348 0.111364854928861 0.124108852481343
## 0                0                1                1
## 1                1                0                0
##      fit
##      0.160270316835722 0.213555588794544 0.241936733479724
## 0                1                1                2
## 1                0                0                0
##      fit
##      0.245766072880757 0.3019444494419796 0.324920443723405
## 0                0                1                0
## 1                1                0                1
##      fit
##      0.342485768826136 0.373135061856799 0.385879059409281
```

```

##      0      0      3      1
##      1      2      0      0
##      fit
##      0.397672202215958 0.413676336263139 0.426420333815621 0.48449991678566
##      0      0      1      0      1
##      1      1      0      1      0
##      fit
##      0.499172092308174 0.537785188744482 0.541369182098166
##      0      1      0      1
##      1      1      1      0
##      fit
##      0.573579666184379 0.586690650651343 0.597792949684554
##      0      0      0      0
##      1      1      1      1
##      fit
##      0.602327797784042 0.614120940590719 0.618901134737556
##      0      0      1      0
##      1      1      1      1
##      fit
##      0.626864938143201 0.642869072190382 0.644770233621383
##      0      0      0      0
##      1      3      2      1
##      fit
##      0.687823553829076 0.726436650265384 0.730265989666417
##      0      0      0      0
##      1      1      1      2
##      fit
##      0.754233927119243 0.758063266520276 0.770807264072758
##      0      0      0      0
##      1      1      1      1
##      fit
##      0.775342112172246 0.799188408757938 0.839729683164278
##      0      1      1      1
##      1      0      0      3
##      fit
##      0.848093870664799 0.886706967101107 0.888635145071139
##      0      0      0      1
##      1      2      1      1
##      fit
##      0.914504243954965 0.946714728041178 0.95945872559366 0.987256002447518
##      0      1      1      1      0
##      1      0      0      1      1
##      fit
##      1
##      0 1
##      1 22

```

```

(fit.train1 <- train.kknn(R1 ~ ., train, kmax = 3,
  kernel = "optimal", distance = 1))

```

```

##
## Call:
## train.kknn(formula = R1 ~ ., data = train, kmax = 3, distance = 1, kernel = "optimal")
##
## Type of response variable: continuous

```

```

## minimal mean absolute error: 0.1991646
## Minimal mean squared error: 0.1514675
## Best kernel: optimal
## Best k: 3

(fit.train2 <- train.kknn(R1 ~ ., train, kmax = 4,
  kernel = "optimal", distance = 1))

##
## Call:
## train.kknn(formula = R1 ~ ., data = train, kmax = 4, distance = 1,      kernel = "optimal")
##
## Type of response variable: continuous
## minimal mean absolute error: 0.1991646
## Minimal mean squared error: 0.1408923
## Best kernel: optimal
## Best k: 4

(fit.train3 <- train.kknn(R1 ~ ., train, kmax = 5,
  kernel = "optimal", distance = 1))

##
## Call:
## train.kknn(formula = R1 ~ ., data = train, kmax = 5, distance = 1,      kernel = "optimal")
##
## Type of response variable: continuous
## minimal mean absolute error: 0.1991646
## Minimal mean squared error: 0.134424
## Best kernel: optimal
## Best k: 5

(fit.train4 <- train.kknn(R1 ~ ., train, kmax = 6,
  kernel = "optimal", distance = 1))

##
## Call:
## train.kknn(formula = R1 ~ ., data = train, kmax = 6, distance = 1,      kernel = "optimal")
##
## Type of response variable: continuous
## minimal mean absolute error: 0.1991646
## Minimal mean squared error: 0.130609
## Best kernel: optimal
## Best k: 6

# table(predict(fit.train1, validate), validate$R1)
# table(predict(fit.train2, validate), validate$R1)
# table(predict(fit.train3, validate), validate$R1)
# table(predict(fit.train4, validate), validate$R1)

# plot(fit.train1)
# plot(fit.train2)
# plot(fit.train3)
# plot(fit.train4)

cat('KKNN accuracy:',sum(fit == test[,11]) / nrow(test),'\n')

## KKNN accuracy: 0.4580153

```